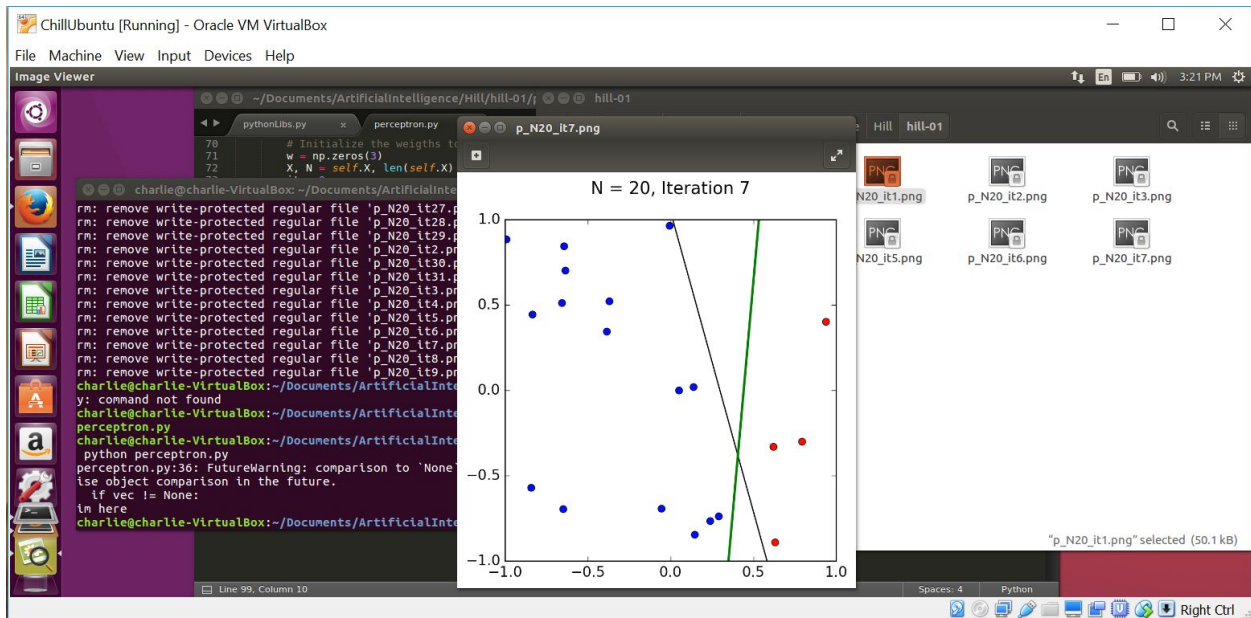


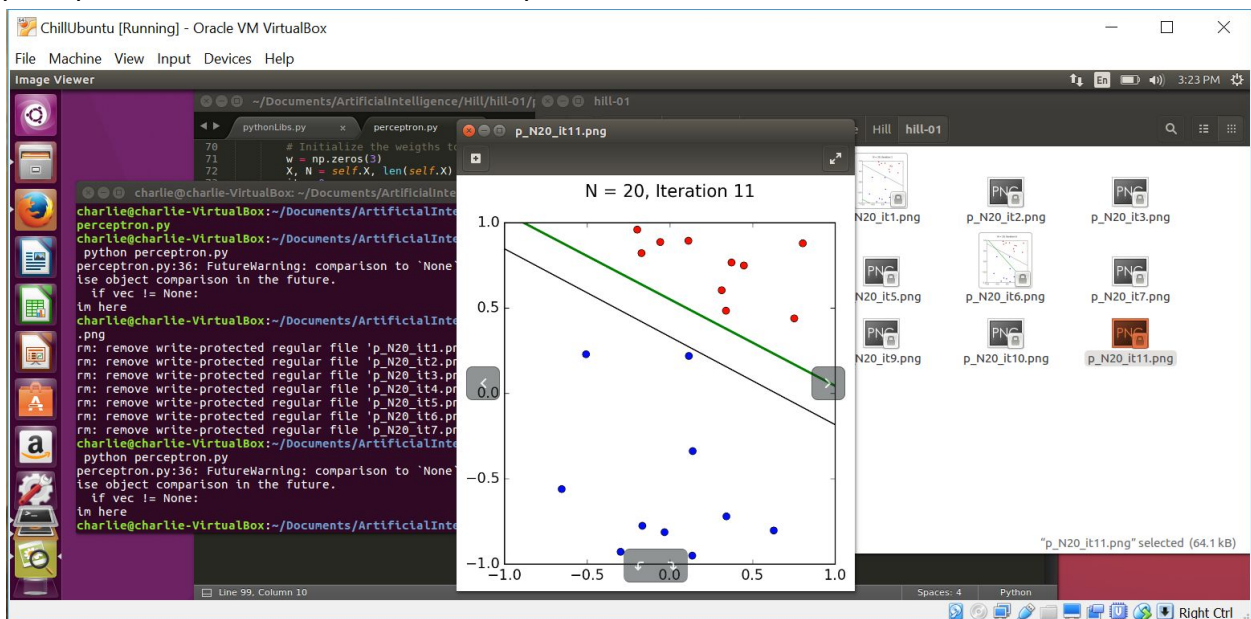
Charlie Hill
Artificial Intelligence
Pablo Rivas
19 September 2016

Homework 01

a and b) Running the perceptron with 20 data points gave me this. It took 7 iterations for the perceptron to find a solution. The perceptron does not seem close to the expected solution that is pre generated.



c) I ran the script again and got this. This time it took 11 iterations to find a solution. The perceptron seems much closer to the expected answer this time.

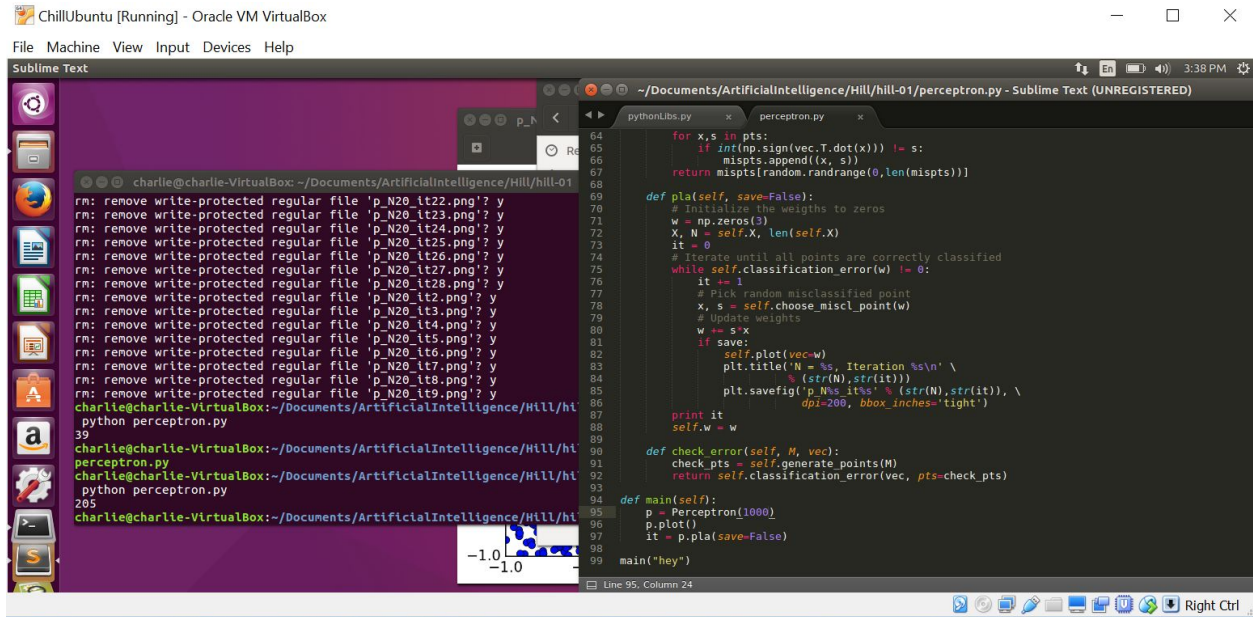


The screenshot displays a Linux desktop environment. On the left is a vertical dock with application icons including a terminal, file manager, web browser, and others. The main workspace contains three windows:

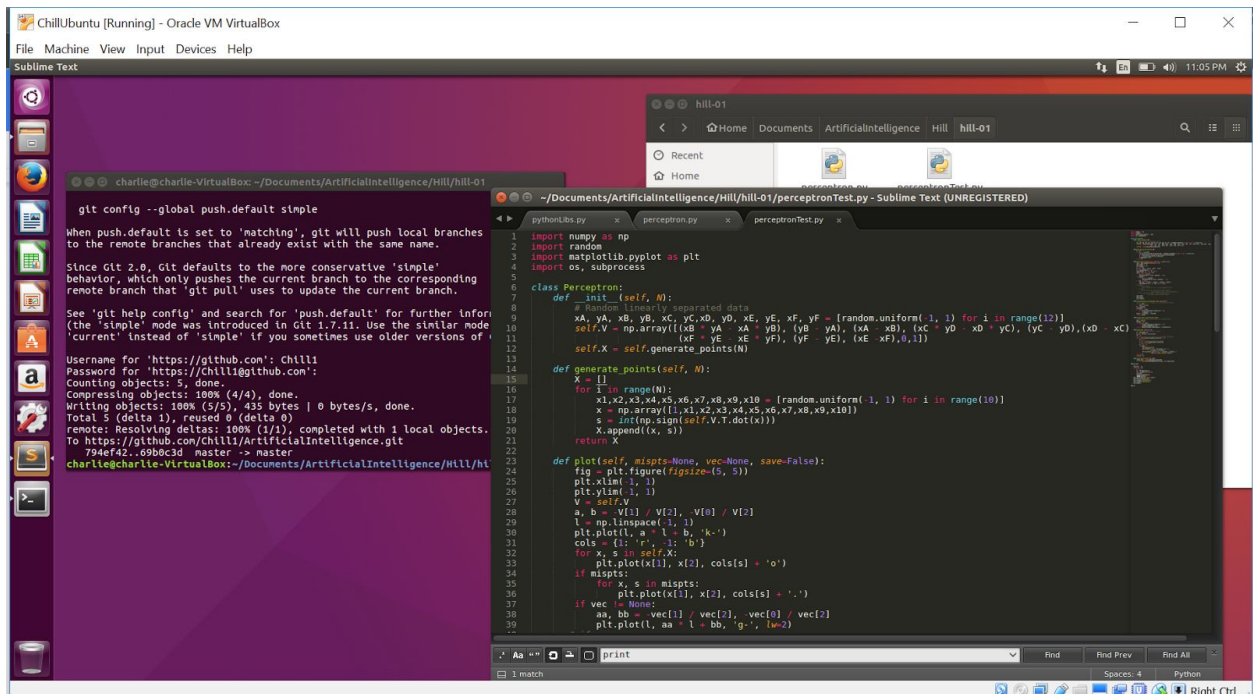
- Terminal Window:** Shows the command prompt at `charlie@charlie-VirtualBox: ~/Documents/ArtificialInte`. It lists several files to be removed (e.g., `p_N20_it10.png`) and then displays a warning message: `python perceptron.py: FutureWarning: comparison to 'None' is object comparison in the future.` followed by a message about memory usage: `max_open_warning, RuntimeWarning) in here`.
- File Explorer Window:** Shows the directory `~/Documents/ArtificialIntelligence/hill/hill-01/`. It contains a grid of image files named `p_N100_it0.png` through `p_N100_it23.png`. The file `p_N100_it24.png` is selected, with a status bar indicating its size is 90.9 kB.
- Plot Window:** A scatter plot titled `N = 100, Iteration 24`. The x-axis ranges from -1.0 to 1.0, and the y-axis ranges from -1.0 to 1.0. The plot shows two classes of data points (blue and red) and a green line representing the decision boundary. The legend indicates that `p_N100_it24.png` is selected.

The screenshot shows a Linux desktop environment with a terminal window, a file manager, and an image viewer. The terminal window displays the execution of a perceptron algorithm, showing the removal of write-protected files and the output of the algorithm. The file manager shows a directory of generated images, including 'p_N1000_it17.png' which is selected. The image viewer shows a scatter plot of data points (blue and red) with a decision boundary line, labeled 'N = 1000, Iteration 17'.

So then I decided to not save the plots so I could find out how many iterations it would actually take. It would end up taking 205 iterations to find a solution. If I had to guess I can imagine the perceptron probably looks really similar to the expected solution because with so many data points the line has to be exact because there is only a little room for solutions. The more we increase the data points the closer the perceptron will get to the expected solution.



f) This was by far the hardest part of the homework to figure out in my opinion. Altering the program to allow for 10 dimensions was not easy. However, I was eventually able to figure it out and get it working. From there running the program became much slower. I imagine the perceptron has much more thinking to do in 10 dimensions.



I ran the program with 1000 data points. It took 2652 iterations to find the perceptron in this screenshot. I ran it a few times and found around 4000 iterations seemed to be pretty common.

as seen by the 1000 data points experiments. As the number of dimensions increase the time it takes for the perceptron to calculate its next move seems to increase. This was noticeable when the 10 dimensions were used to try and calculate the 1000 datapoint perceptron. When there were only two dimensions it seemed to find the solution much faster.